

## Spring 2025 course announcement

# MATH 392: Intro to Neural Networks

Instructor: Arvind Suresh ([arvindsuresh@arizona.edu](mailto:arvindsuresh@arizona.edu))

Credits: 3 (counts toward the 'Application Course' requirement for math majors)

Prereqs: [MATH 223](#)

Can potentially be waived if student has some experience with multi-variable functions or matrices/vectors; **please contact the instructor to enquire!**

(Note: more or less time will be spent as needed to cover the requisite mathematical content, the pre-requisites are mainly to ensure a reasonable pace can be maintained.)

Coding prereqs: Experience with Python is useful but not needed because a lot of code will already be provided, and we will make systematic use of [Github Copilot](#) to get by with minimal coding.

How to sign up: Email instructor to verify pre-requisites (can skip this if you've taken MATH 223), then fill out enrollment form in MATH Rm 108 (main office in the math building). Email [math-academics@arizona.edu](mailto:math-academics@arizona.edu) for more assistance in signing up.

When: Mon-Wed, 11 – 12:15 pm

Where: Modern Languages, Rm 201.

Goal:

To provide students with a self-contained introduction to the mathematics and practical implementation of neural networks, which are a fundamental class of machine learning models that underlie modern AI's like ChatGPT.

Reasons why you might like to enroll:

- You are interested in AI and would like a self-contained introduction covering the basics.
- You are considering internships/jobs in the field of machine learning/AI and would like to get started building a portfolio with projects to showcase to potential employers.
- You enjoy courses that blend theory (from math) with practice (coding).
- You are a math major looking to fill the "Application Course" credit.

Learning Objectives:

- Understand the key mathematical concepts used in neural networks, including linear algebra, gradient descent, and backpropagation.
- Learn to build and implement simple neural networks using libraries like PyTorch.
- Analyze and evaluate neural network models, with an emphasis on model optimization, regularization, and hyperparameter tuning.

- Gain experience in the research method (namely, asking questions and being able to hunt down answers or resources).
- Prepare for independent research by developing the ability to approach problems related to neural networks and machine learning with a solid mathematical framework.

Course features:

- Math concepts will be motivated by asking natural questions about datasets provided by the instructor.
- Every week, students will learn key topics and immediately engage with the material through hands-on coding exercises (Jupyter notebooks prepared before-hand by the instructor).
- Assignments will primarily consist of mini-projects, implemented in Python using industry-standard packages (*sklearn* and *PyTorch*).
- Students will maintain a GitHub repository containing their mini-projects and final project.
- Students will learn to use [GitHub Copilot for Education](#) (free and powerful AI code assistant) to write code with minimal effort.
- Heavy emphasis will be placed on the process of doing research, namely, asking lots of interesting questions and collaborating with peers on projects.

## **Schedule by week**

**(tentative, might spend more time on math/coding portions depending on the background and interests of students)**

Week 1: Introduction to Machine Learning and Neural Networks

### Topics

1. Overview of Machine Learning
  - Introduction to supervised learning.
  - Applications and importance of machine learning.
2. Universal Approximation Theorem
  - Statement and implications for neural networks.
3. Case Study: Image Classification in PyTorch
  - Example of a simple feedforward neural network for image classification.
  - Code demonstration of building, training, and evaluating a model on a small dataset (e.g., MNIST).

### Learning Objectives

- Understand supervised learning as a key machine learning paradigm.
- Appreciate the theoretical significance of the universal approximation theorem.
- Gain initial exposure to neural network implementation in PyTorch.

### Hands-On Example

- Use PyTorch to demonstrate:
  - Loading and visualizing a dataset.
  - Defining a simple network architecture.
  - Training the model.
  - Evaluating and visualizing predictions.

## Week 2: Mathematical Foundations – Linear Algebra (I)

### Topics

- Vectors and Matrices
- Definition and operations: addition, scalar multiplication, and dot products.
- Matrix multiplication and its geometric interpretation.
- Identity matrix and inverse matrix.
- Applications in Neural Networks
- Representing data as vectors (e.g., input features as vectors).
- Matrix multiplication in neural network computations (e.g., inputs and weights).
- The role of the dot product in calculating activations.

### Learning Objectives

- Understand vector and matrix operations.
- Be able to apply matrix multiplication to represent neural network calculations.

### Hands-On Session

- Implement basic matrix operations using NumPy.
- Demonstrate how a neural network layer computes its output using matrix multiplication.

## Week 3: Mathematical Foundations – Basic Statistics

### Topics

- Descriptive Statistics
- Mean, median, and mode.
- Variance, standard deviation, and range.
- Interquartile range and outliers.
- Correlation and Covariance
- Covariance: measuring the relationship between two variables.
- Correlation coefficient: normalizing covariance for easier interpretation.
- Interpretation of correlation in datasets.
- Importance of Data Distribution
- Normal distribution and its significance in data analysis.
- Skewness and kurtosis.

### Learning Objectives

- Understand and calculate key statistics: mean, variance, correlation, etc.
- Learn to interpret statistical measures and their impact on modeling.

### Hands-On Session

- Calculate statistics for a sample dataset using Python (NumPy/Pandas).
- Visualize distributions using histograms and box plots.

## Week 4: Mathematical Foundations – Basic Probability

### Topics

- Introduction to Probability
- Basic probability rules: addition, multiplication, and conditional probability.
- Probability distributions: uniform, normal, and binomial distributions.
- Random variables: discrete vs continuous.
- Bayes' Theorem and Conditional Probability
- Conditional probability and Bayes' theorem.
- Applications in machine learning, especially in classification.
- Expectation and Variance in Probability
- Expected value of a random variable.
- Variance and standard deviation in probability distributions.

### Learning Objectives

- Understand fundamental probability concepts: conditional probability, random variables, and distributions.
- Learn how to apply Bayes' theorem and basic probability rules in data analysis.

### Hands-On Session

- Use Python to simulate random variables and calculate basic probabilities.
- Generate and visualize probability distributions.

## Week 5: Introduction to Loss Functions in Machine Learning

### Topics

1. What is a Loss Function?
  - Definition and role in machine learning.
  - Difference between loss and cost (loss per sample vs average loss).
2. Common Loss Functions (MSE for regression, cross-entropy for classification)
3. Visualizing Loss
  - Loss landscapes and gradients.
  - Intuition behind minimizing loss during training.
4. Connecting Loss to Optimization
  - Gradient descent as a method to minimize loss.
  - Backpropagation as a mechanism to compute gradients in neural networks.

### Learning Objectives

- Understand the purpose of loss functions in machine learning.
- Learn about commonly used loss functions for regression and classification.
- Appreciate the connection between loss functions and optimization.

### Hands-On Session

- Explore different loss functions using synthetic datasets.
- Use Python to compute MSE and Cross-Entropy Loss on example datasets.
- Visualize how changes in predictions affect the loss values.

## Week 6: Linear Regression and Its Connection to Neural Networks

### Topics

- Linear Regression Overview
- Basics of fitting a line to data using least squares.
- Loss function for regression: Mean Squared Error (MSE).
- Multivariable linear regression: extending to multiple input features.
- Link to Neural Networks
- Neural networks as generalizations of linear models.
- Understanding linear layers in neural networks.

### Learning Objectives

- Understand the mechanics of linear regression and how it minimizes loss.
- Recognize how linear models relate to the structure of neural networks.

### Hands-On Session

- Implement linear regression from scratch using Python.
- Use Scikit-learn to quickly fit a linear regression model to a dataset and visualize results.



## Week 7: The Perceptron – Concepts, History, and the XOR Problem

### Topics

1. The Perceptron
  - Concept: a single-layer binary classifier.
  - Mathematical formulation: input, weights, bias, activation function.
  - Training via weight updates: the perceptron learning rule.
2. History of the Perceptron
  - Invented by Frank Rosenblatt in 1958.
  - Early promise and excitement in artificial intelligence.
3. The XOR Problem
  - Demonstration of a linearly non-separable problem.
  - Limitations of the perceptron in solving non-linear problems.
  - Introduction to multi-layer networks as a solution.

### Learning Objectives

- Understand the perceptron as a foundational building block for neural networks.
- Appreciate the historical significance of the perceptron and its limitations.
- Recognize the need for non-linearity and multi-layer networks.

### Hands-On Session

- Implement the perceptron algorithm for a simple linearly separable dataset.
- Visualize the decision boundary for the perceptron.
- Attempt to classify XOR data and discuss why it fails.

## Week 8: Introduction to Multilayer Perceptrons (MLPs)

### Topics

- What is a Multilayer Perceptron?
- Extension of the perceptron with hidden layers.
- Activation functions enabling non-linear transformations.
- Why MLPs Solve the XOR Problem
- Representing non-linear decision boundaries with multiple layers.

### Learning Objectives

- Understand the architecture and function of an MLP.
- Recognize how MLPs overcome the limitations of single-layer perceptrons.

### Hands-On Session

- Implement an MLP for the XOR problem using PyTorch.
- Train and evaluate the model, visualizing the learned decision boundary.

## Week 9: Introduction to Optimization via Gradient Descent

### Topics

- Basics of Gradient Descent
- Concept of gradients and their role in optimization.
- Learning rate and its impact.
- Variants of Gradient Descent
- Batch, Stochastic, and Mini-Batch Gradient Descent.

### Learning Objectives

- Understand how gradient descent is used to optimize loss functions.
- Learn the differences between gradient descent variants.

### Hands-On Session

- Visualize gradient descent on a 2D function.
- Use PyTorch's autograd to implement parameter updates in a simple example.

## Week 10: Backpropagation

### Topics

- Backpropagation: its role in computing gradients and training neural networks.
- Applying backpropagation to single-layer networks.
- Extending backpropagation to multi-layer networks using the chain rule.

### Learning Objectives

- Understand how backpropagation computes gradients layer by layer.
- Learn its application in training multi-layer neural networks.

### Hands-On Session

- Manual implementation of backpropagation for a single-layer network.
- Use PyTorch to observe backpropagation in a multi-layer perceptron.

## Week 11: Activation Functions

### Topics

- The purpose of activation functions: enabling non-linearity in neural networks.
- Common activation functions: Sigmoid, Tanh, ReLU, and their modern variants.
- Choosing the right activation function for a task.

### Learning Objectives

- Understand the importance of activation functions in neural networks.
- Learn the properties, advantages, and limitations of commonly used activation functions.

### Hands-On Session

- Visualize and compare activation functions.
- Experiment with different activation functions in a PyTorch model.

## Week 12: Regularization in Neural Networks

### Topics

- Why Regularization?
- Overfitting and its causes in machine learning models.
- Common Regularization Techniques:
- L1 and L2 regularization (penalizing large weights).
- Dropout: randomly deactivating neurons during training.
- Data augmentation: enhancing the dataset to improve generalization.

### Learning Objectives

- Understand the importance of regularization in training neural networks.
- Learn and apply various regularization techniques.

### Hands-On Session

- Implement L2 regularization and dropout in a PyTorch model.
- Experiment with how regularization impacts model performance on a small dataset.

## Week 13: Model Evaluation and Hyperparameter Tuning

### Topics

- Model Evaluation
- Train, validation, and test splits.
- Common evaluation metrics: accuracy, precision, recall, F1-score.
- Visualizing performance with confusion matrices.
- Hyperparameter Tuning
- Key hyperparameters: learning rate, batch size, number of layers, etc.
- Grid search and random search.
- Introduction to advanced techniques like Bayesian optimization.

### Learning Objectives

- Learn to evaluate models effectively using appropriate metrics.
- Understand the impact of hyperparameters on model training and performance.

### Hands-On Session

- Evaluate a trained model using performance metrics in PyTorch.
- Perform a simple grid search for hyperparameter tuning.